

May 2, 2019

```

In [355]: import keras
import pandas as pd
import numpy as np
from keras import Sequential
import matplotlib.pyplot as plt
from keras.layers import Dense, Conv1D, MaxPooling1D, Flatten, Input, LSTM, Embedding
import tensorflow as tf
from keras.backend.tensorflow_backend import set_session
config = tf.ConfigProto()
config.gpu_options.allow_growth = True
config.log_device_placement = True
sess = tf.Session(config=config)
set_session(sess)
sess.as_default()
sess.graph.as_default()

Out[355]: <contextlib._GeneratorContextManager at 0x7f8e095763c8>

In [25]: data = np.genfromtxt('seqtrain.csv', delimiter=',')
train_X, train_Y = data[:, :29], data[:, 29:]
print(train_X.shape, train_Y.shape)

(4000, 29) (4000, 4)

In [32]: test_data = np.genfromtxt('seqvalid.csv', delimiter=',')
test_X, test_Y = test_data[:, :29], test_data[:, 29:]
print(test_X.shape, test_Y.shape)

(1000, 29) (1000, 4)

In [211]: model = Sequential()
model.add(Dense(units=50, activation='relu', input_dim=29))
model.add(Dense(units=25, activation='relu', input_dim=29))
model.add(Dense(units=10, activation='relu', input_dim=29))
model.add(Dense(units=5, activation='relu', input_dim=29))
model.add(Dense(units=4, activation='softmax'))

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```
In [130]: history = model.fit(train_X, train_Y, epochs=150, batch_size=128, validation_data=(t
```


```
In [131]: model.evaluate(test_X, test_Y)
```

```
1000/1000 [=====] - 0s 63us/step
```

```
Out[131]: [1.8841854248046874, 0.504]
```

```
In [132]: plt.plot(history.history['acc'], label='Train accuracy')
plt.plot(history.history['val_acc'], label='Validation accuracy')
plt.legend()
```

```
Out[132]: <matplotlib.legend.Legend at 0x7f8e61fc2668>
```



```
In [445]: model = Sequential()
model.add(Dense(units=14, activation='relu', input_dim=29))
model.add(Dense(units=7, activation='relu', input_dim=29))
model.add(Dense(units=4, activation='softmax'))


model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(train_X, train_Y, epochs=150, batch_size=128, validation_data=(t
```

```
In [446]: model.summary()
```

Layer (type)	Output Shape	Param #
dense_292 (Dense)	(None, 14)	420
dense_293 (Dense)	(None, 7)	105
dense_294 (Dense)	(None, 4)	32
Total params: 557		
Trainable params: 557		
Non-trainable params: 0		

```
In [435]: plt.plot(history.history['acc'], label='Train accuracy')
          plt.plot(history.history['val_acc'], label='Validation accuracy')
          plt.legend()
```

```
Out[435]: <matplotlib.legend.Legend at 0x7f8dcf2d7eb8>
```



```
In [313]: model.summary()
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

conv1d_72 (Conv1D)	(None, 29, 3)	12

conv1d_73 (Conv1D)	(None, 29, 7)	70

conv1d_74 (Conv1D)	(None, 29, 14)	308

conv1d_75 (Conv1D)	(None, 29, 28)	420

flatten_23 (Flatten)	(None, 812)	0

dense_196 (Dense)	(None, 10)	8130

dense_197 (Dense)	(None, 5)	55

dense_198 (Dense)	(None, 4)	24
=====		
Total params: 9,019		
Trainable params: 9,019		
Non-trainable params: 0		

```
In [336]: model = Sequential()

        model.add(Conv1D(filters=3, kernel_size=3, activation='relu', padding='same', input_shape=(None, 29, 3)))


        model.add(Conv1D(filters=7, kernel_size=5, activation='relu', padding='same'))
        model.add(Conv1D(filters=14, kernel_size=3, activation='relu', padding='same'))
        model.add(Flatten())
        model.add(Dense(10, activation='relu'))
        model.add(Dense(5, activation='relu'))
        model.add(Dense(4, activation='softmax'))

        model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

In [337]: history = model.fit(np.expand_dims(train_X, axis=2), train_Y, epochs=75, batch_size=128)

In [338]: plt.plot(history.history['acc'], label='Train accuracy')
        plt.plot(history.history['val_acc'], label='Train accuracy')
        plt.legend()

Out[338]: <matplotlib.legend.Legend at 0x7f8e0967f898>
```



```
In [339]: model.summary()
```

Layer (type)	Output Shape	Param #
conv1d_96 (Conv1D)	(None, 29, 3)	12
conv1d_97 (Conv1D)	(None, 29, 7)	112
conv1d_98 (Conv1D)	(None, 29, 14)	308
flatten_33 (Flatten)	(None, 406)	0
dense_218 (Dense)	(None, 10)	4070
dense_219 (Dense)	(None, 5)	55
dense_220 (Dense)	(None, 4)	24
Total params: 4,581		
Trainable params: 4,581		
Non-trainable params: 0		

```
In [340]: model = Sequential()
```

```

model.add(Conv1D(filters=3, kernel_size=3, activation='relu', padding='same', input_shape=(None, 290, 1)))
model.add(Conv1D(filters=3, kernel_size=3, activation='relu', padding='same', input_shape=(None, 290, 1)))
model.add(Flatten())
model.add(Dense(5, activation='relu'))
model.add(Dense(4, activation='softmax'))

```

```

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

```

In [334]: history = model.fit(np.expand_dims(train_X, axis=2), train_Y, epochs=50, batch_size=128)

```

```


In [335]: plt.plot(history.history['acc'], label='Train accuracy')
plt.plot(history.history['val_acc'], label='Train accuracy')
plt.legend()

```

```

Out[335]: <matplotlib.legend.Legend at 0x7f8e09e1d518>

```



```

In [385]: model.summary()

```

Layer (type)	Output Shape	Param #
embedding_13 (Embedding)	(None, None, 290)	8410
lstm_41 (LSTM)	(None, None, 128)	214528
lstm_42 (LSTM)	(None, None, 64)	49408

```

-----
dense_242 (Dense)                (None, None, 15)                975
-----
dense_243 (Dense)                (None, None, 4)                 64
=====
Total params: 273,385
Trainable params: 273,385
Non-trainable params: 0
-----

```

```

In [422]: model = Sequential()

          model.add(Embedding(29, output_dim=29))
          model.add(LSTM(29, activation='relu'))
          model.add(Dense(10, activation='relu'))
          model.add(Dense(4, activation='softmax'))


          model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

In [424]: history = model.fit(train_X, train_Y, epochs=150, batch_size=128, verbose=0, validation_data=(val_X, val_Y))

In [425]: plt.plot(history.history['acc'], label='Train accuracy')
          plt.plot(history.history['val_acc'], label='Train accuracy')
          plt.legend()

Out[425]: <matplotlib.legend.Legend at 0x7f8dd04b99e8>

```



```
In [426]: model.summary()
```

Layer (type)	Output Shape	Param #
embedding_23 (Embedding)	(None, None, 29)	841
lstm_52 (LSTM)	(None, 29)	6844
dense_266 (Dense)	(None, 10)	300
dense_267 (Dense)	(None, 4)	44

Total params: 8,029
Trainable params: 8,029
Non-trainable params: 0

```
In [415]: model = Sequential()  
          model.add(Dense(4, activation='softmax'))
```

```
In [416]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
          history = model.fit(train_X, train_Y, epochs=50, batch_size=64, verbose=0, validation_data=(test_X, test_Y))
```

```
In [417]: plt.plot(history.history['acc'], label='Train accuracy')  
          plt.plot(history.history['val_acc'], label='Train accuracy')  
          plt.legend()
```

```
Out[417]: <matplotlib.legend.Legend at 0x7f8dd2179780>
```




```
In [418]: model.summary()
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_263 (Dense)            (None, 4)                 120
=====
Total params: 120
Trainable params: 120
Non-trainable params: 0
-----
```

Fully connected, CNN and RNN all reached pretty much the same about 55% accuracy on the validation set. Adding more layers or neurons really only helps to overfit and the validation accuracy never seems to go over 60% with the few architectures that I tested. Clearly the logistic regression model is really bad and at 25% accuracy it is no better than guessing, so at least a few layers are needed to achieve decent results on this dataset. Generally on this problem a smaller network seems to work best, maybe that is due to the relatively low dimensionality of the problem.

```
In [430]: model = Sequential()
          model.add(Dense(units=500, activation='relu'))
          model.add(Dense(units=250, activation='relu'))
          model.add(Dense(units=100, activation='relu'))
          model.add(Dense(units=50, activation='relu'))
          model.add(Dense(units=4, activation='softmax'))

          model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [432]: history = model.fit(train_X, train_Y, epochs=50, batch_size=64, verbose=1, validation_data=(val_X, val_Y))
```

Train on 4000 samples, validate on 1000 samples

```
Epoch 1/50
4000/4000 [=====] - 0s 89us/step - loss: 0.7201 - acc: 0.6877 - val_loss: 0.6877
Epoch 2/50
4000/4000 [=====] - 0s 90us/step - loss: 0.6490 - acc: 0.7242 - val_loss: 0.6490
Epoch 3/50
4000/4000 [=====] - 0s 88us/step - loss: 0.6126 - acc: 0.7433 - val_loss: 0.6126
Epoch 4/50
4000/4000 [=====] - 0s 90us/step - loss: 0.5657 - acc: 0.7652 - val_loss: 0.5657
Epoch 5/50
4000/4000 [=====] - 0s 89us/step - loss: 0.4823 - acc: 0.8037 - val_loss: 0.4823
Epoch 6/50
4000/4000 [=====] - 0s 91us/step - loss: 0.4137 - acc: 0.8448 - val_loss: 0.4137
Epoch 7/50
4000/4000 [=====] - 0s 88us/step - loss: 0.3443 - acc: 0.8648 - val_loss: 0.3443
```

```

Epoch 8/50
4000/4000 [=====] - 0s 90us/step - loss: 0.2825 - acc: 0.8920 - val_1
Epoch 9/50
4000/4000 [=====] - 0s 90us/step - loss: 0.2262 - acc: 0.9237 - val_1
Epoch 10/50
4000/4000 [=====] - 0s 91us/step - loss: 0.1947 - acc: 0.9345 - val_1
Epoch 11/50
4000/4000 [=====] - 0s 86us/step - loss: 0.1567 - acc: 0.9460 - val_1
Epoch 12/50
4000/4000 [=====] - 0s 84us/step - loss: 0.1440 - acc: 0.9517 - val_1
Epoch 13/50
4000/4000 [=====] - 0s 88us/step - loss: 0.0948 - acc: 0.9725 - val_1
Epoch 14/50
4000/4000 [=====] - 0s 93us/step - loss: 0.0547 - acc: 0.9882 - val_1
Epoch 15/50
4000/4000 [=====] - 0s 85us/step - loss: 0.0315 - acc: 0.9955 - val_1
Epoch 16/50
4000/4000 [=====] - 0s 82us/step - loss: 0.0211 - acc: 0.9982 - val_1
Epoch 17/50
4000/4000 [=====] - 0s 82us/step - loss: 0.0320 - acc: 0.9925 - val_1
Epoch 18/50
4000/4000 [=====] - 0s 88us/step - loss: 0.0781 - acc: 0.9740 - val_1
Epoch 19/50
4000/4000 [=====] - 0s 96us/step - loss: 0.1338 - acc: 0.9527 - val_1
Epoch 20/50
4000/4000 [=====] - 0s 91us/step - loss: 0.1403 - acc: 0.9483 - val_1
Epoch 21/50
4000/4000 [=====] - 0s 85us/step - loss: 0.0969 - acc: 0.9688 - val_1
Epoch 22/50
4000/4000 [=====] - 0s 82us/step - loss: 0.0407 - acc: 0.9888 - val_1
Epoch 23/50
4000/4000 [=====] - 0s 91us/step - loss: 0.0389 - acc: 0.9895 - val_1
Epoch 24/50
4000/4000 [=====] - 0s 91us/step - loss: 0.0174 - acc: 0.9963 - val_1
Epoch 25/50
4000/4000 [=====] - 0s 84us/step - loss: 0.0081 - acc: 0.9988 - val_1
Epoch 26/50
4000/4000 [=====] - 0s 84us/step - loss: 0.0025 - acc: 1.0000 - val_1
Epoch 27/50
4000/4000 [=====] - 0s 83us/step - loss: 0.0014 - acc: 1.0000 - val_1
Epoch 28/50
4000/4000 [=====] - 0s 91us/step - loss: 0.0011 - acc: 1.0000 - val_1
Epoch 29/50
4000/4000 [=====] - 0s 90us/step - loss: 9.0256e-04 - acc: 1.0000 - val_1
Epoch 30/50
4000/4000 [=====] - 0s 87us/step - loss: 7.8576e-04 - acc: 1.0000 - val_1
Epoch 31/50
4000/4000 [=====] - 0s 95us/step - loss: 6.9064e-04 - acc: 1.0000 - val_1

```

```

Epoch 32/50
4000/4000 [=====] - 0s 90us/step - loss: 6.1655e-04 - acc: 1.0000 - va
Epoch 33/50
4000/4000 [=====] - 0s 92us/step - loss: 5.5558e-04 - acc: 1.0000 - va
Epoch 34/50
4000/4000 [=====] - 0s 86us/step - loss: 5.0239e-04 - acc: 1.0000 - va
Epoch 35/50
4000/4000 [=====] - 0s 95us/step - loss: 4.5934e-04 - acc: 1.0000 - va
Epoch 36/50
4000/4000 [=====] - 0s 95us/step - loss: 4.2123e-04 - acc: 1.0000 - va
Epoch 37/50
4000/4000 [=====] - 0s 98us/step - loss: 3.8703e-04 - acc: 1.0000 - va
Epoch 38/50
4000/4000 [=====] - 0s 94us/step - loss: 3.5936e-04 - acc: 1.0000 - va
Epoch 39/50
4000/4000 [=====] - 0s 99us/step - loss: 3.3367e-04 - acc: 1.0000 - va
Epoch 40/50
4000/4000 [=====] - 0s 97us/step - loss: 3.0978e-04 - acc: 1.0000 - va
Epoch 41/50
4000/4000 [=====] - 0s 87us/step - loss: 2.8999e-04 - acc: 1.0000 - va
Epoch 42/50
4000/4000 [=====] - 0s 91us/step - loss: 2.7190e-04 - acc: 1.0000 - va
Epoch 43/50
4000/4000 [=====] - 0s 89us/step - loss: 2.5476e-04 - acc: 1.0000 - va
Epoch 44/50
4000/4000 [=====] - 0s 89us/step - loss: 2.3919e-04 - acc: 1.0000 - va
Epoch 45/50
4000/4000 [=====] - 0s 90us/step - loss: 2.2538e-04 - acc: 1.0000 - va
Epoch 46/50
4000/4000 [=====] - 0s 99us/step - loss: 2.1189e-04 - acc: 1.0000 - va
Epoch 47/50
4000/4000 [=====] - 0s 89us/step - loss: 1.9991e-04 - acc: 1.0000 - va
Epoch 48/50
4000/4000 [=====] - 0s 96us/step - loss: 1.8886e-04 - acc: 1.0000 - va
Epoch 49/50
4000/4000 [=====] - 0s 95us/step - loss: 1.7833e-04 - acc: 1.0000 - va
Epoch 50/50
4000/4000 [=====] - 0s 94us/step - loss: 1.6866e-04 - acc: 1.0000 - va

```

```


In [433]: plt.plot(history.history['acc'], label='Train accuracy')
          plt.plot(history.history['val_acc'], label='Train accuracy')
          plt.legend()

```

```

Out[433]: <matplotlib.legend.Legend at 0x7f8dcf7917f0>

```



In []:

In []:

In []: